




# Pattern Recognition



# Speech Emotion Recognition

Alaa Hossam Abu-hashima	6750
Mohammad Helaly	6870
Mohammad Shamarka	6952

Dr/Marwan Torki

Eng/Abdelrahman Wael

# Speech Emotion Recognition

## Importing Libraries

```
! pip install audiomentations
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: audiomentations in /usr/local/lib/python3.10/dist-packages (0.30.0)
Requirement already satisfied: numpy>=1.13.0 in /usr/local/lib/python3.10/dist-packages (from audiomentations) (1.22.4)
Requirement already satisfied: librosa<0.10.0,>0.7.2 in /usr/local/lib/python3.10/dist-packages (from audiomentations) (0.9.2)
Requirement already satisfied: scipy<2,>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from audiomentations) (1.10.1)
Requirement already satisfied: audioread>=2.1.9 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (3.0.1)
Requirement already satisfied: scikit-learn>=0.19.1 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (1.2.2)
Requirement already satisfied: joblib>=0.14 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (1.3.2)
Requirement already satisfied: decorator>=4.0.10 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (4.4.2)
Requirement already satisfied: resampy>=0.2.2 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (0.4.2)
Requirement already satisfied: numba>=0.45.1 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (0.56.0)
Requirement already satisfied: soundfile>=0.10.2 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (0.12.1)
Requirement already satisfied: pooch>=1.0 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (1.6.0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from librosa<0.10.0,>0.7.2->audiomentations) (23.1)
Requirement already satisfied: llvmlite<0.40,>=0.39.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba>=0.45.1->librosa<0.10.0,>0.7.2->audiomentations) (0.40.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from numba>=0.45.1->librosa<0.10.0,>0.7.2->audiomentations) (68.0.0)
Requirement already satisfied: appdirs>=1.3.0 in /usr/local/lib/python3.10/dist-packages (from pooch>=1.0->librosa<0.10.0,>0.7.2->audiomentations) (1.4.4)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from pooch>=1.0->librosa<0.10.0,>0.7.2->audiomentations) (2.31.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.19.1->librosa<0.10.0,>0.7.2->audiomentations) (3.1.0)
Requirement already satisfied: cffi>=1.0 in /usr/local/lib/python3.10/dist-packages (from soundfile>=0.10.2->librosa<0.10.0,>0.7.2->audiomentations) (1.16.0)
Requirement already satisfied: pycparser in /usr/local/lib/python3.10/dist-packages (from cffi>=1.0->soundfile>=0.10.2->librosa<0.10.0,>0.7.2->audiomentations) (2.21)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->pooch>=1.0->librosa<0.10.0,>0.7.2->audiomentations) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->pooch>=1.0->librosa<0.10.0,>0.7.2->audiomentations) (2023.7.22)
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->pooch>=1.0->librosa<0.10.0,>0.7.2->audiomentations) (2.0.12)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->pooch>=1.0->librosa<0.10.0,>0.7.2->audiomentations) (3.4)

```

```
import numpy as np
import pandas as pd
import os.path
from scipy.io import wavfile
import librosa
import soundfile
import seaborn as sns
```

```
import matplotlib.pyplot as plt
import IPython.display as ipd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler, MinMaxScaler
from sklearn.metrics import confusion_matrix, f1_score, accuracy_score
from audiomentations import Compose, AddGaussianNoise, TimeStretch, PitchShift, Shift, Gain, SpecCompose, SpecChannelShuffle, SpecFrequencyMask
import tensorflow as tf
from keras import utils
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Conv1D, MaxPooling1D, Conv2D, MaxPooling2D, Flatten, BatchNormalization
from keras import optimizers, regularizers
```

## Preprocessing Data

### Loading Dataset

```
#kaggle datasets download -d dmitrybabko/speech-emotion-recognition-en
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

### Listing the audios files in array (only thier names)

```
#path = '../input/speech-emotion-recognition-en/Crema'
path = '/content/drive/MyDrive/Crema'

files = []

for file in os.listdir(path):
    if file.endswith('.wav'):
        files.append(os.path.join(path, file))

#print(len(files))
#for i in range(5):
#    print(files[i])
```

```
#print("\n")

labels = []

for file in files:
    label = os.path.splitext(os.path.split(file)[1])[0].split('_')[2]
    labels.append(label)

print(len(labels))
print(set(labels))
dict_label={"HAP": "Happy", "ANG": "Anger", "SAD": "Sad", "NEU": "Neutral", "DIS": "Disgust", "FEA": "Fear" }
```

```
7451
{'HAP', 'NEU', 'SAD', 'DIS', 'ANG', 'FEA'}
```

loading the raw data of each audio in array

```
audios = []

for i in files:
    audio, sr = librosa.load(i, sr= None)
    audios.append(audio)

print(len(audios))
print(len(audios[0]))
for i in range(5):
    print(audios[i])
```

```
7451
36303
[-1.3916016e-02 -1.5411377e-02 -1.4587402e-02 ... -9.1552734e-05
 3.0517578e-05 -6.1035156e-05]
[0.00567627 0.00515747 0.00448608 ... 0.          0.          0.          ]
[-0.00027466 -0.00018311 0.00024414 ... 0.          0.          0.          ]
[ 0.00100708 0.00045776 -0.00033569 ... 0.          0.          0.          ]
[ 0.00024414 -0.00039673 -0.00045776 ... 0.          0.          0.          ]
```

Audio padding

To overcome the problem that the audios are not the same length, we get the max audio length in the data and make padding for all the audios that have less length

```
max_length = max(len(audio) for audio in audios)

padded_audios = []
for audio in audios:
    padding = max_length - len(audio)
    padded_audio = np.pad(audio, (0, padding), mode='constant')
    padded_audios.append(padded_audio)

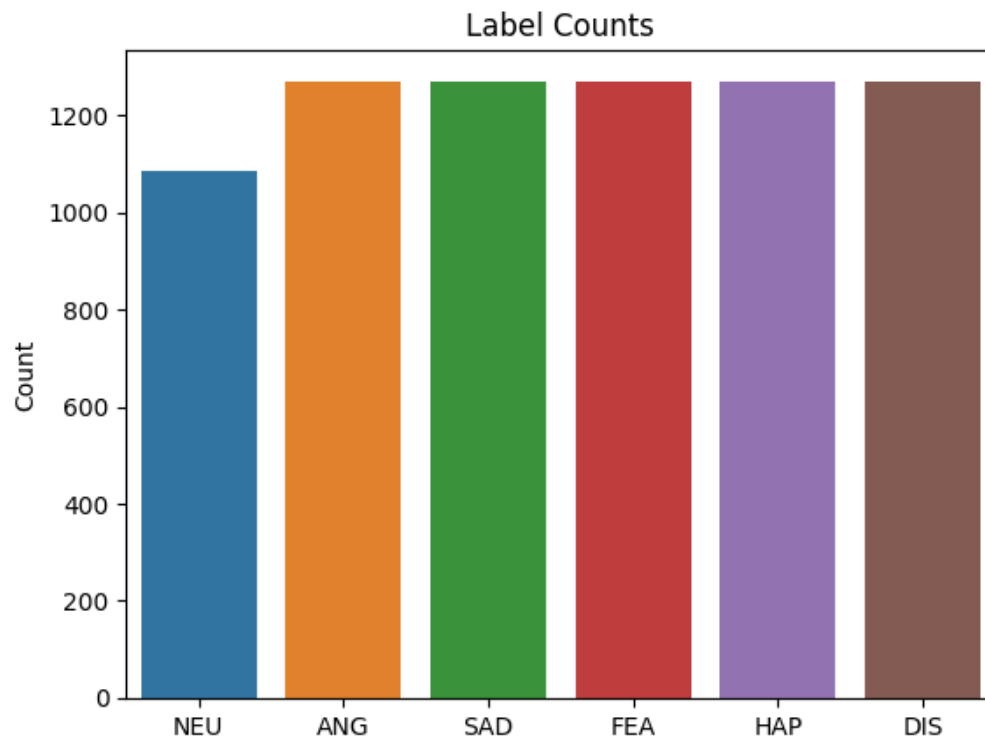
print(len(padded_audios))
print(len(padded_audios[0]))
for i in range(5):
    print(padded_audios[i])
```

7451  
80080  
[-0.01391602 -0.01541138 -0.0145874 ... 0. 0.  
0. ]  
[0.00567627 0.00515747 0.00448608 ... 0. 0. 0. ]  
[-0.00027466 -0.00018311 0.00024414 ... 0. 0.  
0. ]  
[ 0.00100708 0.00045776 -0.00033569 ... 0. 0.  
0. ]  
[ 0.00024414 -0.00039673 -0.00045776 ... 0. 0.  
0. ]

## Plotting Label Counts

This plot shows the number of audios of each label.

```
sns.countplot(x=labels)
plt.title("Label Counts")
plt.xlabel("Label")
plt.ylabel("Count")
plt.show()
```




## Plotting Raw Signal of Different Audios

This function plots the spectrum of the different types of audios.


```
def plot_audio(start):  
    for label in dict_label.keys():  
        audio_num=labels.index(label,start)  
        fig=plt.figure()  
        plt.plot(audios[audio_num])  
        fig.suptitle(str(dict_label[labels[audio_num]]+" Audio"))  
        plt.xlabel("t")  
        plt.ylabel("f(t)")  
        plt.show()
```

```
plot_audio(0)
```

Happy Audio



Anger Audio





## Playing Audios

This function makes you able to display the audio according to the label(happiness, anger, fear, neutral , sad, disgust and so on) and the index of the audio.

```
def play_audio(label,x):  
    i=labels.index(label,x)  
    print(dict_label[label]+" Audio")  
    return ipd.Audio(files[i])
```



```
play_audio("HAP",20)
```

Happy Audio



```
play_audio("DIS",10)
```

Disgust Audio



```
play_audio("ANG",30)
```

Anger Audio




```
play_audio("SAD",60)
```

Sad Audio









play\_audio("FEA",60)

Fear Audio




play\_audio("SAD",50)

Sad Audio




play\_audio("HAP",60)

Happy Audio




play\_audio("FEA",70)

Fear Audio



play\_audio("NEU",10)

Neutral Audio



play\_audio("ANG",90)

Anger Audi



```
play_audio("ANG", 900)
```

Anger Audio



## Encoding Labels

Anger



Anger

Encoding the categorical data (labels).

Anger

```
y = np.array(labels)
y = y.reshape(-1, 1)
```

```
encoder = OneHotEncoder(sparse_output=False)
encoder.fit(y)
```

```
y = encoder.transform(y)
#print(set(y))
print(y.shape)
```

(7451, 6)

## Splitting Data

Splitting the data into train and test sets.

```
X_train_val, X_test, y_train_val, y_test = train_test_split(np.asarray(padded_audios), y, test_size=0.3, random_state=42, stratify = y)
X_train, X_val, y_train, y_val = train_test_split(X_train_val, y_train_val, test_size=0.05, random_state=42, stratify = y_train_val)
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
print(X_val.shape)
print(y_val.shape)
```

```
(4954, 80080)
(4954, 6)
(2236, 80080)
(2236, 6)
(261, 80080)
(261, 6)
```

## Data Augmentation

This part to be used in 2D model to enhance the accuracy and training of the model but colab ram is crashed so we didn't use it unfortunately.

```
augmentation = Compose([AddGaussianNoise(p=0.5),
                        TimeStretch(p=0.5),
                        PitchShift(p=0.5),
                        Shift(p=0.5)])
```

```
augmented_audios = []
for audio in X_train:
    augmented_audio = augmentation(samples=audio, sample_rate=sr)
    augmented_audios.append(augmented_audio)
augmented_audios = np.array(augmented_audios)
print(augmented_audios.shape)
```

```
y_train = np.repeat(y_train, len(augmented_audios))
X_train = np.concatenate((padded_audios, augmented_audios), axis=0)

print(X_train.shape)
print(y.shape)
```

## Extracting Features

### Frequency Domain

Extracting features from frequency domain.

## MFCCs

```
train_mfcc = []

for audio in X_train:
    mfccs = np.mean(librosa.feature.mfcc(y=audio),axis= 0)
    train_mfcc.append(mfccs)

#train_mfcc = np.squeeze(np.array(train_mfcc), axis=1)
train_mfcc = np.expand_dims(train_mfcc, -1)
print(train_mfcc.shape)

test_mfcc = []

for audio in X_test:
    mfccs = np.mean(librosa.feature.mfcc(y=audio),axis= 0)
    test_mfcc.append(mfccs)

#test_mfcc = np.squeeze(np.array(test_mfcc), axis=1)
test_mfcc = np.expand_dims(test_mfcc, -1)
print(test_mfcc.shape)

val_mfcc = []

for audio in X_val:
    mfccs = np.mean(librosa.feature.mfcc(y=audio),axis= 0)
    val_mfcc.append(mfccs)

#val_mfcc = np.squeeze(np.array(val_mfcc), axis=1)
val_mfcc = np.expand_dims(val_mfcc, -1)
print(val_mfcc.shape)

(4948, 157, 1)
(2233, 157, 1)
(261, 157, 1)
```

## Spectral Roll-off

```
train_sr = []

for audio in X_train:
```

```

    srf = librosa.feature.spectral_rolloff(y=audio)
    train_sr.append(srf)

train_sr = np.squeeze(np.array(train_sr), axis=1)
train_sr = np.expand_dims(train_sr, -1)
print(train_sr.shape)

test_sr = []

for audio in X_test:
    srf = librosa.feature.spectral_rolloff(y=audio)
    test_sr.append(srf)

test_sr = np.squeeze(np.array(test_sr), axis=1)
test_sr = np.expand_dims(test_sr, -1)
print(test_sr.shape)

val_sr = []

for audio in X_val:
    srf = librosa.feature.spectral_rolloff(y=audio)
    val_sr.append(srf)

val_sr = np.squeeze(np.array(val_sr), axis=1)
val_sr = np.expand_dims(val_sr, -1)
print(val_sr.shape)

(4948, 157, 1)
(2233, 157, 1)
(261, 157, 1)

```

## Time Domain

Extracting features from Time domain.

### Energy

```

train_energy = []

for audio in X_train:

```

```

    energy = librosa.feature.rms(y=audio)
    train_energy.append(energy)

train_energy = np.squeeze(np.array(train_energy), axis=1)
train_energy = np.expand_dims(train_energy, -1)
print(train_energy.shape)

test_energy = []

for audio in X_test:
    energy = librosa.feature.rms(y=audio)
    test_energy.append(energy)

test_energy = np.squeeze(np.array(test_energy), axis=1)
test_energy = np.expand_dims(test_energy, -1)
print(test_energy.shape)

val_energy = []

for audio in X_val:
    energy = librosa.feature.rms(y=audio)
    val_energy.append(energy)

val_energy = np.squeeze(np.array(val_energy), axis=1)
val_energy = np.expand_dims(val_energy, -1)
print(val_energy.shape)

(4954, 157, 1)
(2236, 157, 1)
(261, 157, 1)

```

## Zero Crossing Rate

```

train_zcr = []

for audio in X_train:
    zcr = librosa.feature.zero_crossing_rate(audio)
    train_zcr.append(zcr)

train_zcr = np.squeeze(np.array(train_zcr), axis=1)
train_zcr = np.expand_dims(train_zcr, -1)
print(train_zcr.shape)

```

```
test_zcr = []

for audio in X_test:
    zcr = librosa.feature.zero_crossing_rate(audio)
    test_zcr.append(zcr)

test_zcr = np.squeeze(np.array(test_zcr), axis=1)
test_zcr = np.expand_dims(test_zcr, -1)
print(test_zcr.shape)

val_zcr = []

for audio in X_val:
    zcr = librosa.feature.zero_crossing_rate(audio)
    val_zcr.append(zcr)

val_zcr = np.squeeze(np.array(val_zcr), axis=1)
val_zcr = np.expand_dims(val_zcr, -1)
print(val_zcr.shape)

(4954, 157, 1)
(2236, 157, 1)
(261, 157, 1)
```

## Mel Spectrogram

```
train_mel = []

for audio in X_train:
    mel_spec = librosa.feature.melspectrogram(y=audio, sr=sr)
    mel_spec = librosa.power_to_db(mel_spec, ref=np.max)
    train_mel.append(mel_spec)

train_mel = np.array(train_mel)
train_mel = np.expand_dims(train_mel, -1)
print(train_mel.shape)

test_mel = []

for audio in X_test:
    mel_spec = librosa.feature.melspectrogram(y=audio, sr=sr)
    mel_spec = librosa.power_to_db(mel_spec, ref=np.max)
```

```

    test_mel.append(mel_spec)

test_mel = np.array(test_mel)
test_mel = np.expand_dims(test_mel, -1)
print(test_mel.shape)

val_mel = []

for audio in X_val:
    mel_spec = librosa.feature.melspectrogram(y=audio, sr=sr)
    mel_spec = librosa.power_to_db(mel_spec, ref=np.max)
    val_mel.append(mel_spec)

val_mel = np.array(val_mel)
val_mel = np.expand_dims(val_mel, -1)
print(val_mel.shape)

(4954, 128, 157, 1)
(2236, 128, 157, 1)
(261, 128, 157, 1)

```


Plotting of spectrogram.

```

plt.imshow(mel_spec, aspect='auto', origin='lower', cmap='jet')
plt.colorbar(format='%+2.0f dB')
plt.xlabel('Time')
plt.ylabel('Mel Frequency')
plt.show()

```





## CNN Models



## Model 1D - Energy



Testing 1D model with energy feature only.



## Building Model

```
model = Sequential()
model.add(Conv1D(16, 12, activation='relu', input_shape=(train_energy.shape[1], 1)))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 8, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 5, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
#model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(6, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.001),
              metrics=['accuracy'])

model.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
=====		
conv1d_12 (Conv1D)	(None, 146, 16)	208
max_pooling1d_12 (MaxPooling1D)	(None, 73, 16)	0
conv1d_13 (Conv1D)	(None, 66, 24)	3096
max_pooling1d_13 (MaxPooling1D)	(None, 33, 24)	0
conv1d_14 (Conv1D)	(None, 29, 24)	2904
max_pooling1d_14 (MaxPooling1D)	(None, 14, 24)	0
flatten_4 (Flatten)	(None, 336)	0
dense_8 (Dense)	(None, 64)	21568
dropout_4 (Dropout)	(None, 64)	0
dense_9 (Dense)	(None, 6)	390
=====		
Total params: 28,166		
Trainable params: 28,166		
Non-trainable params: 0		

```
history = model.fit(train_energy, y_train,
                    epochs=50,
                    batch_size=128,
                    validation_data=(val_energy, y_val))
```

```
Epoch 1/50
39/39 [=====] - 2s 28ms/step - loss: 1.7254 - accuracy: 0.2084 - val_loss: 1.6094 - val_accuracy: 0.2874
Epoch 2/50
39/39 [=====] - 1s 23ms/step - loss: 1.5927 - accuracy: 0.3145 - val_loss: 1.4802 - val_accuracy: 0.3870
Epoch 3/50
39/39 [=====] - 1s 22ms/step - loss: 1.5160 - accuracy: 0.3608 - val_loss: 1.4089 - val_accuracy: 0.4444
Epoch 4/50
39/39 [=====] - 1s 24ms/step - loss: 1.4870 - accuracy: 0.3895 - val_loss: 1.3773 - val_accuracy: 0.4751
Epoch 5/50
```

39/39 [=====] - 1s 22ms/step - loss: 1.4643 - accuracy: 0.4000 - val\_loss: 1.3642 - val\_accuracy: 0.4828  
Epoch 6/50  
39/39 [=====] - 1s 22ms/step - loss: 1.4571 - accuracy: 0.4052 - val\_loss: 1.3585 - val\_accuracy: 0.4713  
Epoch 7/50  
39/39 [=====] - 1s 21ms/step - loss: 1.4477 - accuracy: 0.4084 - val\_loss: 1.3491 - val\_accuracy: 0.4943  
Epoch 8/50  
39/39 [=====] - 1s 22ms/step - loss: 1.4473 - accuracy: 0.4139 - val\_loss: 1.3467 - val\_accuracy: 0.4636  
Epoch 9/50  
39/39 [=====] - 1s 22ms/step - loss: 1.4364 - accuracy: 0.4165 - val\_loss: 1.3451 - val\_accuracy: 0.4904  
Epoch 10/50  
39/39 [=====] - 1s 21ms/step - loss: 1.4296 - accuracy: 0.4115 - val\_loss: 1.3239 - val\_accuracy: 0.5019  
Epoch 11/50  
39/39 [=====] - 1s 21ms/step - loss: 1.4227 - accuracy: 0.4204 - val\_loss: 1.3145 - val\_accuracy: 0.4904  
Epoch 12/50  
39/39 [=====] - 1s 20ms/step - loss: 1.4177 - accuracy: 0.4210 - val\_loss: 1.3118 - val\_accuracy: 0.5134  
Epoch 13/50  
39/39 [=====] - 1s 20ms/step - loss: 1.4158 - accuracy: 0.4274 - val\_loss: 1.3099 - val\_accuracy: 0.5019  
Epoch 14/50  
39/39 [=====] - 1s 19ms/step - loss: 1.4085 - accuracy: 0.4333 - val\_loss: 1.3049 - val\_accuracy: 0.5057  
Epoch 15/50  
39/39 [=====] - 1s 20ms/step - loss: 1.4074 - accuracy: 0.4335 - val\_loss: 1.2946 - val\_accuracy: 0.5096  
Epoch 16/50  
39/39 [=====] - 1s 20ms/step - loss: 1.4074 - accuracy: 0.4272 - val\_loss: 1.2987 - val\_accuracy: 0.5096  
Epoch 17/50  
39/39 [=====] - 1s 22ms/step - loss: 1.3986 - accuracy: 0.4317 - val\_loss: 1.3333 - val\_accuracy: 0.4789  
Epoch 18/50  
39/39 [=====] - 1s 23ms/step - loss: 1.3952 - accuracy: 0.4319 - val\_loss: 1.3007 - val\_accuracy: 0.5172  
Epoch 19/50  
39/39 [=====] - 1s 22ms/step - loss: 1.4048 - accuracy: 0.4307 - val\_loss: 1.2958 - val\_accuracy: 0.5096  
Epoch 20/50  
39/39 [=====] - 1s 21ms/step - loss: 1.3944 - accuracy: 0.4313 - val\_loss: 1.3080 - val\_accuracy: 0.5057  
Epoch 21/50  
39/39 [=====] - 1s 23ms/step - loss: 1.3928 - accuracy: 0.4400 - val\_loss: 1.2761 - val\_accuracy: 0.5057  
Epoch 22/50  
39/39 [=====] - 1s 23ms/step - loss: 1.3869 - accuracy: 0.4418 - val\_loss: 1.2910 - val\_accuracy: 0.5019  
Epoch 23/50  
39/39 [=====] - 1s 22ms/step - loss: 1.3861 - accuracy: 0.4456 - val\_loss: 1.2755 - val\_accuracy: 0.5172  
Epoch 24/50  
39/39 [=====] - 1s 22ms/step - loss: 1.3833 - accuracy: 0.4402 - val\_loss: 1.2925 - val\_accuracy: 0.5096  
Epoch 25/50  
39/39 [=====] - 1s 22ms/step - loss: 1.3889 - accuracy: 0.4357 - val\_loss: 1.2690 - val\_accuracy: 0.5364  
Epoch 26/50  
39/39 [=====] - 1s 21ms/step - loss: 1.3753 - accuracy: 0.4448 - val\_loss: 1.2580 - val\_accuracy: 0.5326  
Epoch 27/50  
39/39 [=====] - 1s 22ms/step - loss: 1.3780 - accuracy: 0.4438 - val\_loss: 1.2675 - val\_accuracy: 0.5172  
Epoch 28/50  
39/39 [=====] - 1s 20ms/step - loss: 1.3695 - accuracy: 0.4410 - val\_loss: 1.2669 - val\_accuracy: 0.5211  
Epoch 29/50

## Evaluating Model

```
y_pred = model.predict(test_energy)
print("\n")
print("Accuracy score =", accuracy_score(y_test.argmax(axis=1), y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'macro'))
print("\n")
print("F1 score (weighted average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'weighted'))
```


70/70 [=====] - 0s 3ms/step

Accuracy score = 0.44872369010300045

F1 score (macro average) = 0.436351352454696

F1 score (weighted average) = 0.43693261772795944

```
con_mat = confusion_matrix(y_test.argmax(axis=1), y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                 xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                 yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```
class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)
#print(precisions)
#print(recalls)

class_idxes = list(range(len(class_labels)))
sorted_classes = sorted(class_idxes, key=lambda x: f1_scores[x])

k=1
print("Most Confusing Classes:\n")
```

```

for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

1 - FEA:	F1-score = 0.2835	Precision = 0.3543	Recall = 0.2362
2 - DIS:	F1-score = 0.3442	Precision = 0.4530	Recall = 0.2775
3 - NEU:	F1-score = 0.4124	Precision = 0.4421	Recall = 0.3865
4 - HAP:	F1-score = 0.4413	Precision = 0.3792	Recall = 0.5276
5 - SAD:	F1-score = 0.5113	Precision = 0.4502	Recall = 0.5916
6 - ANG:	F1-score = 0.6255	Precision = 0.5911	Recall = 0.6640

## Model 1D - Zero Crossing Rate

Testing 1D model with Zero crossing rate feature only.

### Building Model

```

model = Sequential()
model.add(Conv1D(16, 12, activation='relu', input_shape=(train_zcr.shape[1], 1)))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 8, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 5, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
#model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(6, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.001),

```

```
metrics=['accuracy'])
```

```
model.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====		
conv1d_15 (Conv1D)	(None, 146, 16)	208
max_pooling1d_15 (MaxPooling1D)	(None, 73, 16)	0
conv1d_16 (Conv1D)	(None, 66, 24)	3096
max_pooling1d_16 (MaxPooling1D)	(None, 33, 24)	0
conv1d_17 (Conv1D)	(None, 29, 24)	2904
max_pooling1d_17 (MaxPooling1D)	(None, 14, 24)	0
flatten_5 (Flatten)	(None, 336)	0
dense_10 (Dense)	(None, 64)	21568
dropout_5 (Dropout)	(None, 64)	0
dense_11 (Dense)	(None, 6)	390
=====		
Total params: 28,166		
Trainable params: 28,166		
Non-trainable params: 0		

```
history = model.fit(train_zcr, y_train,  
                    epochs=50,  
                    batch_size=128,  
                    validation_data=(val_zcr, y_val))
```

Epoch 1/50

39/39 [=====] - 2s 29ms/step - loss: 1.7773 - accuracy: 0.1789 - val\_loss: 1.7373 - val\_accuracy: 0.1954

Epoch 2/50

39/39 [=====] - 1s 22ms/step - loss: 1.7441 - accuracy: 0.2591 - val\_loss: 1.6983 - val\_accuracy: 0.2414

Epoch 3/50

39/39 [=====] - 1s 20ms/step - loss: 1.7107 - accuracy: 0.2801 - val\_loss: 1.6455 - val\_accuracy: 0.3142  
Epoch 4/50  
39/39 [=====] - 1s 20ms/step - loss: 1.6579 - accuracy: 0.3173 - val\_loss: 1.5946 - val\_accuracy: 0.3678  
Epoch 5/50  
39/39 [=====] - 1s 19ms/step - loss: 1.6157 - accuracy: 0.3426 - val\_loss: 1.5881 - val\_accuracy: 0.3525  
Epoch 6/50  
39/39 [=====] - 1s 21ms/step - loss: 1.5925 - accuracy: 0.3482 - val\_loss: 1.5510 - val\_accuracy: 0.3563  
Epoch 7/50  
39/39 [=====] - 1s 21ms/step - loss: 1.5815 - accuracy: 0.3498 - val\_loss: 1.5439 - val\_accuracy: 0.3755  
Epoch 8/50  
39/39 [=====] - 1s 21ms/step - loss: 1.5773 - accuracy: 0.3549 - val\_loss: 1.5665 - val\_accuracy: 0.3372  
Epoch 9/50  
39/39 [=====] - 1s 22ms/step - loss: 1.5576 - accuracy: 0.3630 - val\_loss: 1.5819 - val\_accuracy: 0.3180  
Epoch 10/50  
39/39 [=====] - 1s 21ms/step - loss: 1.5706 - accuracy: 0.3517 - val\_loss: 1.5419 - val\_accuracy: 0.3793  
Epoch 11/50  
39/39 [=====] - 1s 22ms/step - loss: 1.5427 - accuracy: 0.3723 - val\_loss: 1.5312 - val\_accuracy: 0.3716  
Epoch 12/50  
39/39 [=====] - 1s 22ms/step - loss: 1.5341 - accuracy: 0.3739 - val\_loss: 1.5386 - val\_accuracy: 0.3716  
Epoch 13/50  
39/39 [=====] - 1s 20ms/step - loss: 1.5366 - accuracy: 0.3763 - val\_loss: 1.5258 - val\_accuracy: 0.3525  
Epoch 14/50  
39/39 [=====] - 1s 21ms/step - loss: 1.5384 - accuracy: 0.3630 - val\_loss: 1.5199 - val\_accuracy: 0.3755  
Epoch 15/50  
39/39 [=====] - 1s 22ms/step - loss: 1.5200 - accuracy: 0.3842 - val\_loss: 1.5329 - val\_accuracy: 0.3678  
Epoch 16/50  
39/39 [=====] - 1s 22ms/step - loss: 1.5172 - accuracy: 0.3820 - val\_loss: 1.5161 - val\_accuracy: 0.3870  
Epoch 17/50  
39/39 [=====] - 1s 22ms/step - loss: 1.5132 - accuracy: 0.3836 - val\_loss: 1.5185 - val\_accuracy: 0.3716  
Epoch 18/50  
39/39 [=====] - 1s 22ms/step - loss: 1.5116 - accuracy: 0.3874 - val\_loss: 1.5256 - val\_accuracy: 0.3525  
Epoch 19/50  
39/39 [=====] - 1s 21ms/step - loss: 1.5108 - accuracy: 0.3858 - val\_loss: 1.5185 - val\_accuracy: 0.3946  
Epoch 20/50  
39/39 [=====] - 1s 20ms/step - loss: 1.4993 - accuracy: 0.3937 - val\_loss: 1.5269 - val\_accuracy: 0.3678  
Epoch 21/50  
39/39 [=====] - 1s 20ms/step - loss: 1.5038 - accuracy: 0.3846 - val\_loss: 1.5316 - val\_accuracy: 0.3678  
Epoch 22/50  
39/39 [=====] - 1s 24ms/step - loss: 1.4944 - accuracy: 0.3981 - val\_loss: 1.5175 - val\_accuracy: 0.3678  
Epoch 23/50  
39/39 [=====] - 1s 25ms/step - loss: 1.4959 - accuracy: 0.3959 - val\_loss: 1.5312 - val\_accuracy: 0.3410  
Epoch 24/50  
39/39 [=====] - 1s 23ms/step - loss: 1.4948 - accuracy: 0.3909 - val\_loss: 1.5277 - val\_accuracy: 0.3257  
Epoch 25/50  
39/39 [=====] - 1s 21ms/step - loss: 1.5028 - accuracy: 0.3899 - val\_loss: 1.5062 - val\_accuracy: 0.3755  
Epoch 26/50  
39/39 [=====] - 1s 20ms/step - loss: 1.4930 - accuracy: 0.3981 - val\_loss: 1.5280 - val\_accuracy: 0.3257  
Epoch 27/50



```

39/39 [=====] - 1s 20ms/step - loss: 1.4873 - accuracy: 0.3961 - val_loss: 1.5223 - val_accuracy: 0.3640
Epoch 28/50
39/39 [=====] - 1s 20ms/step - loss: 1.4810 - accuracy: 0.3973 - val_loss: 1.5163 - val_accuracy: 0.3678
Epoch 29/50
39/39 [=====] - 1s 20ms/step - loss: 1.4766 - accuracy: 0.4004 - val_loss: 1.5100 - val_accuracy: 0.3662

```

## Evaluating Model

```

y_pred = model.predict(test_zcr)
print("\n")
print("Accuracy score =",accuracy_score(y_test.argmax(axis=1),y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =",f1_score(y_test.argmax(axis=1),y_pred.argmax(axis=1),average = 'macro'))
print("\n")
print("F1 score (weighted average) =",f1_score(y_test.argmax(axis=1),y_pred.argmax(axis=1),average = 'weighted'))

```

```

70/70 [=====] - 0s 3ms/step

```

```

Accuracy score = 0.39632781012091356

```

```

F1 score (macro average) = 0.38286335100423435

```

```


F1 score (weighted average) = 0.38328540496165664

```

```

con_mat = confusion_matrix(y_test.argmax(axis=1),y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                 xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                 yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()

```



```

class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)
#print(precisions)
#print(recalls)

class_idxes = list(range(len(class_labels)))
sorted_classes = sorted(class_idxes, key=lambda x: f1_scores[x])

```

```

k=1
print("Most Confusing Classes:\n")
for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

1 - FEA:	F1-score = 0.1735	Precision = 0.2663	Recall = 0.1286
2 - HAP:	F1-score = 0.3488	Precision = 0.3341	Recall = 0.3648
3 - NEU:	F1-score = 0.3675	Precision = 0.3385	Recall = 0.4018
4 - DIS:	F1-score = 0.3746	Precision = 0.4167	Recall = 0.3403
5 - SAD:	F1-score = 0.4864	Precision = 0.4156	Recall = 0.5864
6 - ANG:	F1-score = 0.5464	Precision = 0.5367	Recall = 0.5564

## Model 1D - Energy and Zero Crossing Rate

This model gave us the best results, with accuracy reaching 48%.

### Building Model

```

model = Sequential()
model.add(Conv1D(16, 12, activation='relu', input_shape=((np.concatenate((train_energy, train_zcr), axis=1)).shape[1], 1)))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 8, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 5, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
#model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(6, activation='softmax'))

```

```
model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.001),
              metrics=['accuracy'])
```

```
model.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
=====		
conv1d_18 (Conv1D)	(None, 303, 16)	208
max_pooling1d_18 (MaxPooling1D)	(None, 151, 16)	0
conv1d_19 (Conv1D)	(None, 144, 24)	3096
max_pooling1d_19 (MaxPooling1D)	(None, 72, 24)	0
conv1d_20 (Conv1D)	(None, 68, 24)	2904
max_pooling1d_20 (MaxPooling1D)	(None, 34, 24)	0
flatten_6 (Flatten)	(None, 816)	0
dense_12 (Dense)	(None, 64)	52288
dropout_6 (Dropout)	(None, 64)	0
dense_13 (Dense)	(None, 6)	390
=====		
Total params: 58,886		
Trainable params: 58,886		
Non-trainable params: 0		
=====		

```
history = model.fit(np.concatenate((train_energy, train_zcr), axis=1), y_train,
                    epochs=50,
                    batch_size=128,
                    validation_data=(np.concatenate((val_energy, val_zcr), axis=1), y_val))
```

Epoch 1/50  
39/39 [=====] - 2s 37ms/step - loss: 1.7090 - accuracy: 0.2387 - val\_loss: 1.5530 - val\_accuracy: 0.3372  
Epoch 2/50  
39/39 [=====] - 1s 31ms/step - loss: 1.5649 - accuracy: 0.3482 - val\_loss: 1.4812 - val\_accuracy: 0.3946  
Epoch 3/50  
39/39 [=====] - 1s 33ms/step - loss: 1.5159 - accuracy: 0.3804 - val\_loss: 1.3936 - val\_accuracy: 0.4789  
Epoch 4/50  
39/39 [=====] - 1s 33ms/step - loss: 1.4814 - accuracy: 0.3931 - val\_loss: 1.3811 - val\_accuracy: 0.4100  
Epoch 5/50  
39/39 [=====] - 1s 35ms/step - loss: 1.4526 - accuracy: 0.4095 - val\_loss: 1.3580 - val\_accuracy: 0.4751  
Epoch 6/50  
39/39 [=====] - 1s 36ms/step - loss: 1.4331 - accuracy: 0.4163 - val\_loss: 1.3452 - val\_accuracy: 0.4598  
Epoch 7/50  
39/39 [=====] - 1s 35ms/step - loss: 1.4296 - accuracy: 0.4224 - val\_loss: 1.3220 - val\_accuracy: 0.4674  
Epoch 8/50  
39/39 [=====] - 1s 34ms/step - loss: 1.4150 - accuracy: 0.4220 - val\_loss: 1.3241 - val\_accuracy: 0.4674  
Epoch 9/50  
39/39 [=====] - 1s 35ms/step - loss: 1.3992 - accuracy: 0.4349 - val\_loss: 1.3076 - val\_accuracy: 0.4674  
Epoch 10/50  
39/39 [=====] - 1s 37ms/step - loss: 1.3881 - accuracy: 0.4353 - val\_loss: 1.3304 - val\_accuracy: 0.4559  
Epoch 11/50  
39/39 [=====] - 1s 36ms/step - loss: 1.3812 - accuracy: 0.4509 - val\_loss: 1.3078 - val\_accuracy: 0.4713  
Epoch 12/50  
39/39 [=====] - 1s 31ms/step - loss: 1.3650 - accuracy: 0.4527 - val\_loss: 1.2955 - val\_accuracy: 0.4943  
Epoch 13/50  
39/39 [=====] - 1s 34ms/step - loss: 1.3619 - accuracy: 0.4513 - val\_loss: 1.2965 - val\_accuracy: 0.4713  
Epoch 14/50  
39/39 [=====] - 1s 35ms/step - loss: 1.3521 - accuracy: 0.4543 - val\_loss: 1.2925 - val\_accuracy: 0.4751  
Epoch 15/50  
39/39 [=====] - 1s 31ms/step - loss: 1.3455 - accuracy: 0.4596 - val\_loss: 1.2779 - val\_accuracy: 0.4789  
Epoch 16/50  
39/39 [=====] - 1s 31ms/step - loss: 1.3502 - accuracy: 0.4559 - val\_loss: 1.2726 - val\_accuracy: 0.4828  
Epoch 17/50  
39/39 [=====] - 1s 31ms/step - loss: 1.3366 - accuracy: 0.4598 - val\_loss: 1.2646 - val\_accuracy: 0.4789  
Epoch 18/50  
39/39 [=====] - 1s 33ms/step - loss: 1.3265 - accuracy: 0.4626 - val\_loss: 1.2694 - val\_accuracy: 0.4713  
Epoch 19/50  
39/39 [=====] - 1s 31ms/step - loss: 1.3149 - accuracy: 0.4685 - val\_loss: 1.2591 - val\_accuracy: 0.4943  
Epoch 20/50  
39/39 [=====] - 1s 31ms/step - loss: 1.3152 - accuracy: 0.4802 - val\_loss: 1.2680 - val\_accuracy: 0.4828  
Epoch 21/50  
39/39 [=====] - 1s 32ms/step - loss: 1.3111 - accuracy: 0.4764 - val\_loss: 1.2713 - val\_accuracy: 0.4943  
Epoch 22/50  
39/39 [=====] - 1s 35ms/step - loss: 1.3085 - accuracy: 0.4691 - val\_loss: 1.2734 - val\_accuracy: 0.4866  
Epoch 23/50  
39/39 [=====] - 1s 33ms/step - loss: 1.2910 - accuracy: 0.4830 - val\_loss: 1.2598 - val\_accuracy: 0.4866  
Epoch 24/50  
39/39 [=====] - 1s 33ms/step - loss: 1.2875 - accuracy: 0.4810 - val\_loss: 1.2622 - val\_accuracy: 0.5019

```
Epoch 25/50
39/39 [=====] - 1s 32ms/step - loss: 1.2850 - accuracy: 0.4893 - val_loss: 1.2797 - val_accuracy: 0.4789
Epoch 26/50
39/39 [=====] - 1s 31ms/step - loss: 1.2867 - accuracy: 0.4822 - val_loss: 1.2581 - val_accuracy: 0.5057
Epoch 27/50
39/39 [=====] - 1s 33ms/step - loss: 1.2827 - accuracy: 0.4877 - val_loss: 1.2608 - val_accuracy: 0.4981
Epoch 28/50
39/39 [=====] - 1s 32ms/step - loss: 1.2781 - accuracy: 0.4879 - val_loss: 1.2532 - val_accuracy: 0.4981
Epoch 29/50
39/39 [=====] - 1s 35ms/step - loss: 1.2806 - accuracy: 0.4919 - val_loss: 1.2702 - val_accuracy: 0.4866
```

## Evaluating Model

```
y_pred = model.predict(np.concatenate((test_energy, test_zcr), axis=1))
print("\n")
print("Accuracy score =", accuracy_score(y_test.argmax(axis=1), y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'macro'))
print("\n")
print("F1 score (weighted average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'weighted'))
```

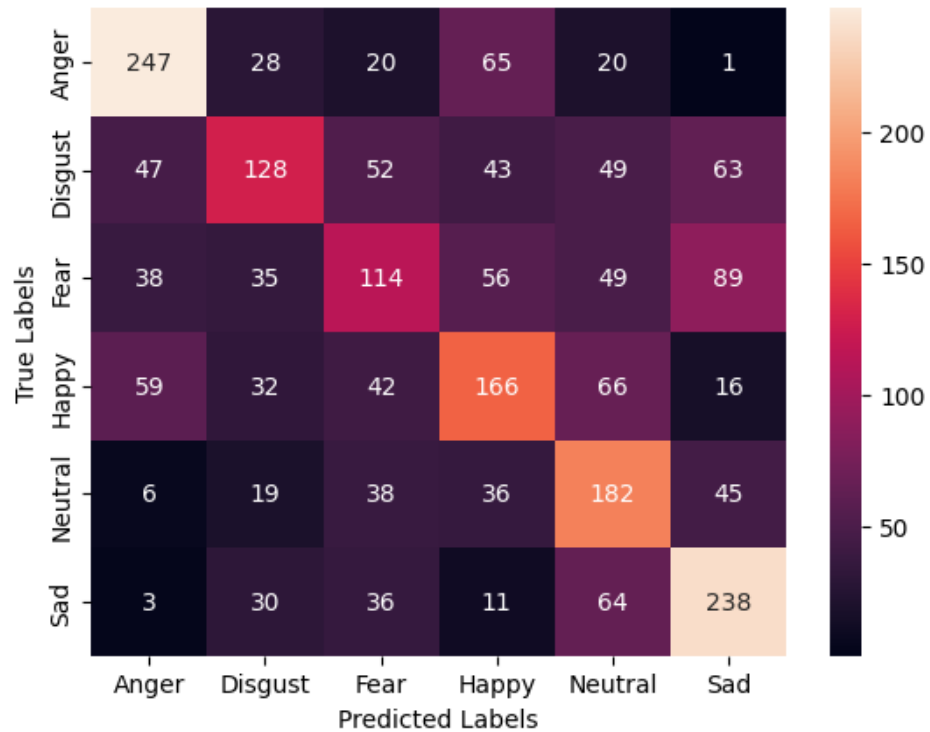
```
70/70 [=====] - 0s 6ms/step
```

```
Accuracy score = 0.48141513658755036
```

```
F1 score (macro average) = 0.47466678367205245
```

```
F1 score (weighted average) = 0.4745046872848074
```

```
con_mat = confusion_matrix(y_test.argmax(axis=1), y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                 xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                 yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```

class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)

```

```

#print(precisions)
#print(recalls)

class_idx = list(range(len(class_labels)))
sorted_classes = sorted(class_idx, key=lambda x: f1_scores[x])

k=1
print("Most Confusing Classes:\n")
for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

1 - FEA:	F1-score = 0.3338	Precision = 0.3775	Recall = 0.2992
2 - DIS:	F1-score = 0.3914	Precision = 0.4706	Recall = 0.3351
3 - HAP:	F1-score = 0.4380	Precision = 0.4403	Recall = 0.4357
4 - NEU:	F1-score = 0.4815	Precision = 0.4233	Recall = 0.5583
5 - SAD:	F1-score = 0.5707	Precision = 0.5265	Recall = 0.6230
6 - ANG:	F1-score = 0.6325	Precision = 0.6175	Recall = 0.6483

## Another 1D model

```

model = Sequential()
model.add(Conv1D(256, 12, activation='relu', input_shape=((np.concatenate((train_energy, train_zcr), axis=1)).shape[1], 1)))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(128, 8, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(64, 5, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Flatten())
#model.add(Dense(128, activation='relu'))
#model.add(Dropout(0.3))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.5))
model.add(BatchNormalization())

```



```

model.add(Dense(6, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.0001),
              metrics=['accuracy'])

model.summary()

```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
=====		
conv1d_21 (Conv1D)	(None, 303, 256)	3328
max_pooling1d_21 (MaxPooling1D)	(None, 151, 256)	0
conv1d_22 (Conv1D)	(None, 144, 128)	262272
max_pooling1d_22 (MaxPooling1D)	(None, 72, 128)	0
conv1d_23 (Conv1D)	(None, 68, 64)	41024
max_pooling1d_23 (MaxPooling1D)	(None, 34, 64)	0
flatten_7 (Flatten)	(None, 2176)	0
dense_18 (Dense)	(None, 64)	139328
dropout_11 (Dropout)	(None, 64)	0
batch_normalization_5 (Batch Normalization)	(None, 64)	256
dense_19 (Dense)	(None, 6)	390
=====		
Total params: 446,598		
Trainable params: 446,470		
Non-trainable params: 128		
=====		

```

history = model.fit(np.concatenate((train_energy, train_zcr), axis=1), y_train,
                    epochs=25,

```

```
batch_size=128,  
validation_data=(np.concatenate((val_energy, val_zcr), axis=1), y_val))
```

```
Epoch 1/25  
39/39 [=====] - 46s 1s/step - loss: 1.7234 - accuracy: 0.2662 - val_loss: 1.7800 - val_accuracy: 0.3525  
Epoch 2/25  
39/39 [=====] - 44s 1s/step - loss: 1.5988 - accuracy: 0.3516 - val_loss: 1.7661 - val_accuracy: 0.3908  
Epoch 3/25  
39/39 [=====] - 45s 1s/step - loss: 1.5539 - accuracy: 0.3692 - val_loss: 1.7529 - val_accuracy: 0.4368  
Epoch 4/25  
39/39 [=====] - 44s 1s/step - loss: 1.5163 - accuracy: 0.3815 - val_loss: 1.7410 - val_accuracy: 0.4100  
Epoch 5/25  
39/39 [=====] - 44s 1s/step - loss: 1.5032 - accuracy: 0.3958 - val_loss: 1.7293 - val_accuracy: 0.3908  
Epoch 6/25  
39/39 [=====] - 45s 1s/step - loss: 1.4869 - accuracy: 0.4025 - val_loss: 1.7148 - val_accuracy: 0.4330  
Epoch 7/25  
39/39 [=====] - 45s 1s/step - loss: 1.4628 - accuracy: 0.4150 - val_loss: 1.6949 - val_accuracy: 0.4368  
Epoch 8/25  
39/39 [=====] - 44s 1s/step - loss: 1.4578 - accuracy: 0.4142 - val_loss: 1.6663 - val_accuracy: 0.4330  
Epoch 9/25  
39/39 [=====] - 44s 1s/step - loss: 1.4490 - accuracy: 0.4235 - val_loss: 1.6403 - val_accuracy: 0.4521  
Epoch 10/25  
39/39 [=====] - 44s 1s/step - loss: 1.4445 - accuracy: 0.4312 - val_loss: 1.6016 - val_accuracy: 0.4444  
Epoch 11/25  
39/39 [=====] - 45s 1s/step - loss: 1.4236 - accuracy: 0.4378 - val_loss: 1.5699 - val_accuracy: 0.4483  
Epoch 12/25  
39/39 [=====] - 45s 1s/step - loss: 1.4216 - accuracy: 0.4374 - val_loss: 1.5470 - val_accuracy: 0.4444  
Epoch 13/25  
39/39 [=====] - 44s 1s/step - loss: 1.4237 - accuracy: 0.4419 - val_loss: 1.5230 - val_accuracy: 0.4483  
Epoch 14/25  
39/39 [=====] - 44s 1s/step - loss: 1.4016 - accuracy: 0.4497 - val_loss: 1.4814 - val_accuracy: 0.4713  
Epoch 15/25  
39/39 [=====] - 44s 1s/step - loss: 1.4043 - accuracy: 0.4514 - val_loss: 1.4647 - val_accuracy: 0.4521  
Epoch 16/25  
39/39 [=====] - 45s 1s/step - loss: 1.3877 - accuracy: 0.4528 - val_loss: 1.4502 - val_accuracy: 0.4521  
Epoch 17/25  
39/39 [=====] - 44s 1s/step - loss: 1.3866 - accuracy: 0.4520 - val_loss: 1.4205 - val_accuracy: 0.4559  
Epoch 18/25  
39/39 [=====] - 44s 1s/step - loss: 1.3834 - accuracy: 0.4621 - val_loss: 1.3999 - val_accuracy: 0.4598  
Epoch 19/25  
39/39 [=====] - 45s 1s/step - loss: 1.3750 - accuracy: 0.4621 - val_loss: 1.4250 - val_accuracy: 0.4483  
Epoch 20/25  
39/39 [=====] - 44s 1s/step - loss: 1.3779 - accuracy: 0.4574 - val_loss: 1.3968 - val_accuracy: 0.4674  
Epoch 21/25  
39/39 [=====] - 44s 1s/step - loss: 1.3657 - accuracy: 0.4647 - val_loss: 1.3949 - val_accuracy: 0.4330  
Epoch 22/25  
39/39 [=====] - 44s 1s/step - loss: 1.3661 - accuracy: 0.4651 - val_loss: 1.3750 - val_accuracy: 0.4866  
Epoch 23/25
```

```
39/39 [=====] - 44s 1s/step - loss: 1.3486 - accuracy: 0.4711 - val_loss: 1.3917 - val_accuracy: 0.4368
Epoch 24/25
39/39 [=====] - 45s 1s/step - loss: 1.3542 - accuracy: 0.4653 - val_loss: 1.3639 - val_accuracy: 0.4751
Epoch 25/25
39/39 [=====] - 44s 1s/step - loss: 1.3443 - accuracy: 0.4717 - val_loss: 1.3925 - val_accuracy: 0.4674
```

```
y_pred = model.predict(np.concatenate((test_energy, test_zcr), axis=1))
print("\n")
print("Accuracy score =",accuracy_score(y_test.argmax(axis=1),y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =",f1_score(y_test.argmax(axis=1),y_pred.argmax(axis=1),average = 'macro'))
print("\n")
print("F1 score (weighted average) =",f1_score(y_test.argmax(axis=1),y_pred.argmax(axis=1),average = 'weighted'))
```


```
70/70 [=====] - 4s 57ms/step
```

```
Accuracy score = 0.4601967799642218
```

```
F1 score (macro average) = 0.4343275560587625
```

```
F1 score (weighted average) = 0.4348376787576417
```

```
con_mat = confusion_matrix(y_test.argmax(axis=1),y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                 xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                 yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```

class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)
#print(precisions)
#print(recalls)

class_idxes = list(range(len(class_labels)))
sorted_classes = sorted(class_idxes, key=lambda x: f1_scores[x])

k=1

```

```

print("Most Confusing Classes:\n")
for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

1 - FEA:	F1-score = 0.2830	Precision = 0.3603	Recall = 0.2330
2 - DIS:	F1-score = 0.4060	Precision = 0.4770	Recall = 0.3534
3 - HAP:	F1-score = 0.4222	Precision = 0.3995	Recall = 0.4476
4 - NEU:	F1-score = 0.4255	Precision = 0.3756	Recall = 0.4908
5 - SAD:	F1-score = 0.5378	Precision = 0.4967	Recall = 0.5864
6 - ANG:	F1-score = 0.6539	Precision = 0.6384	Recall = 0.6702

## Model 1D - MFCCs

### Building Model

```

model = Sequential()
model.add(Conv1D(16, 12, activation='relu', input_shape=(train_mfcc.shape[1], 1)))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 8, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 5, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
#model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(6, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.001),
              metrics=['accuracy'])

```

```
model.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
=====		
conv1d_21 (Conv1D)	(None, 146, 16)	208
max_pooling1d_21 (MaxPooling1D)	(None, 73, 16)	0
conv1d_22 (Conv1D)	(None, 66, 24)	3096
max_pooling1d_22 (MaxPooling1D)	(None, 33, 24)	0
conv1d_23 (Conv1D)	(None, 29, 24)	2904
max_pooling1d_23 (MaxPooling1D)	(None, 14, 24)	0
flatten_7 (Flatten)	(None, 336)	0
dense_14 (Dense)	(None, 64)	21568
dropout_7 (Dropout)	(None, 64)	0
dense_15 (Dense)	(None, 6)	390
=====		
Total params: 28,166		
Trainable params: 28,166		
Non-trainable params: 0		

```
history = model.fit(train_mfcc, y_train,  
                    epochs=50,  
                    batch_size=128,  
                    validation_data=(val_mfcc, y_val))
```

Эпока 25/50

39/39 [=====] - 1s 21ms/step - loss: 1.3595 - accuracy: 0.4450 - val\_loss: 1.4107 - val\_accuracy: 0.3985

Epoch 26/50

39/39 [=====] - 1s 21ms/step - loss: 1.3636 - accuracy: 0.4590 - val\_loss: 1.4106 - val\_accuracy: 0.4138

Epoch 27/50

39/39 [=====] - 1s 21ms/step - loss: 1.3354 - accuracy: 0.4576 - val\_loss: 1.4167 - val\_accuracy: 0.3831

```
Epoch 49/50
39/39 [=====] - 1s 22ms/step - loss: 1.1249 - accuracy: 0.5558 - val_loss: 1.5392 - val_accuracy: 0.3793
Epoch 50/50
39/39 [=====] - 1s 22ms/step - loss: 1.1191 - accuracy: 0.5521 - val_loss: 1.5481 - val_accuracy: 0.3563
```

## Evaluating Model

```
y_pred = model.predict(test_mfcc)
print("\n")
print("Accuracy score =", accuracy_score(y_test.argmax(axis=1), y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'macro'))
print("\n")
print("F1 score (weighted average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'weighted'))
```

```
70/70 [=====] - 0s 3ms/step
```


```
Accuracy score = 0.37214509628302733
```

```
F1 score (macro average) = 0.365114083038715
```

```
F1 score (weighted average) = 0.3663917421448286
```

```
con_mat = confusion_matrix(y_test.argmax(axis=1), y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                 xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                 yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```





```
class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)
#print(precisions)
#print(recalls)

class_idxes = list(range(len(class_labels)))
sorted_classes = sorted(class_idxes, key=lambda x: f1_scores[x])
```

```

k=1
print("Most Confusing Classes:\n")
for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

1 - FEA:	F1-score = 0.1865	Precision = 0.2407	Recall = 0.1522
2 - NEU:	F1-score = 0.3140	Precision = 0.3121	Recall = 0.3160
3 - DIS:	F1-score = 0.3199	Precision = 0.3558	Recall = 0.2906
4 - HAP:	F1-score = 0.3640	Precision = 0.3126	Recall = 0.4357
5 - SAD:	F1-score = 0.4534	Precision = 0.3969	Recall = 0.5288
6 - ANG:	F1-score = 0.5528	Precision = 0.6161	Recall = 0.5013

## Model 1D - Spectral Roll-off

This model gave us the worst results by far.

## Building Model

```

model = Sequential()
model.add(Conv1D(16, 12, activation='relu', input_shape=(train_sr.shape[1], 1)))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 8, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 5, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
#model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(6, activation='softmax'))

```

```
model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.001),
              metrics=['accuracy'])
```

```
model.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
=====		
conv1d_24 (Conv1D)	(None, 146, 16)	208
max_pooling1d_24 (MaxPooling1D)	(None, 73, 16)	0
conv1d_25 (Conv1D)	(None, 66, 24)	3096
max_pooling1d_25 (MaxPooling1D)	(None, 33, 24)	0
conv1d_26 (Conv1D)	(None, 29, 24)	2904
max_pooling1d_26 (MaxPooling1D)	(None, 14, 24)	0
flatten_8 (Flatten)	(None, 336)	0
dense_16 (Dense)	(None, 64)	21568
dropout_8 (Dropout)	(None, 64)	0
dense_17 (Dense)	(None, 6)	390
=====		
Total params: 28,166		
Trainable params: 28,166		
Non-trainable params: 0		
=====		

```
history = model.fit(train_sr, y_train,
                    epochs=50,
                    batch_size=128,
                    validation_data=(val_sr, y_val))
```

Epoch 1/50

39/39 [=====] - 2s 27ms/step - loss: 179.7276 - accuracy: 0.1692 - val\_loss: 19.6667 - val\_accuracy: 0.2184

Epoch 2/50  
39/39 [=====] - 1s 21ms/step - loss: 14.7119 - accuracy: 0.1742 - val\_loss: 6.2626 - val\_accuracy: 0.1877  
Epoch 3/50  
39/39 [=====] - 1s 22ms/step - loss: 5.3804 - accuracy: 0.1684 - val\_loss: 3.2832 - val\_accuracy: 0.1916  
Epoch 4/50  
39/39 [=====] - 1s 22ms/step - loss: 2.7428 - accuracy: 0.1908 - val\_loss: 2.2282 - val\_accuracy: 0.1686  
Epoch 5/50  
39/39 [=====] - 1s 21ms/step - loss: 2.0660 - accuracy: 0.1865 - val\_loss: 1.9669 - val\_accuracy: 0.1648  
Epoch 6/50  
39/39 [=====] - 1s 22ms/step - loss: 1.9271 - accuracy: 0.1880 - val\_loss: 1.9352 - val\_accuracy: 0.1648  
Epoch 7/50  
39/39 [=====] - 1s 22ms/step - loss: 1.8584 - accuracy: 0.1878 - val\_loss: 1.9356 - val\_accuracy: 0.1609  
Epoch 8/50  
39/39 [=====] - 1s 21ms/step - loss: 1.8285 - accuracy: 0.1896 - val\_loss: 1.9182 - val\_accuracy: 0.1648  
Epoch 9/50  
39/39 [=====] - 1s 19ms/step - loss: 1.8152 - accuracy: 0.1916 - val\_loss: 1.9090 - val\_accuracy: 0.1609  
Epoch 10/50  
39/39 [=====] - 1s 21ms/step - loss: 1.7972 - accuracy: 0.1908 - val\_loss: 1.8916 - val\_accuracy: 0.1648  
Epoch 11/50  
39/39 [=====] - 1s 20ms/step - loss: 1.7890 - accuracy: 0.1948 - val\_loss: 1.9022 - val\_accuracy: 0.1609  
Epoch 12/50  
39/39 [=====] - 1s 20ms/step - loss: 1.7785 - accuracy: 0.1956 - val\_loss: 1.8865 - val\_accuracy: 0.1648  
Epoch 13/50  
39/39 [=====] - 1s 19ms/step - loss: 1.7695 - accuracy: 0.1979 - val\_loss: 1.8975 - val\_accuracy: 0.1648  
Epoch 14/50  
39/39 [=====] - 1s 22ms/step - loss: 1.7675 - accuracy: 0.1954 - val\_loss: 1.8973 - val\_accuracy: 0.1609  
Epoch 15/50  
39/39 [=====] - 1s 22ms/step - loss: 1.7673 - accuracy: 0.1962 - val\_loss: 1.8914 - val\_accuracy: 0.1648  
Epoch 16/50  
39/39 [=====] - 1s 21ms/step - loss: 1.7541 - accuracy: 0.1989 - val\_loss: 1.8919 - val\_accuracy: 0.1648  
Epoch 17/50  
39/39 [=====] - 1s 21ms/step - loss: 1.7573 - accuracy: 0.2013 - val\_loss: 1.8996 - val\_accuracy: 0.1571  
Epoch 18/50  
39/39 [=====] - 1s 19ms/step - loss: 1.7554 - accuracy: 0.1975 - val\_loss: 1.8978 - val\_accuracy: 0.1571  
Epoch 19/50  
39/39 [=====] - 1s 22ms/step - loss: 1.7499 - accuracy: 0.2033 - val\_loss: 1.9188 - val\_accuracy: 0.1609  
Epoch 20/50  
39/39 [=====] - 1s 20ms/step - loss: 1.7512 - accuracy: 0.1979 - val\_loss: 1.8866 - val\_accuracy: 0.1877  
Epoch 21/50  
39/39 [=====] - 1s 20ms/step - loss: 1.7461 - accuracy: 0.1999 - val\_loss: 1.9324 - val\_accuracy: 0.1724  
Epoch 22/50  
39/39 [=====] - 1s 22ms/step - loss: 1.7504 - accuracy: 0.1964 - val\_loss: 1.8910 - val\_accuracy: 0.1686  
Epoch 23/50  
39/39 [=====] - 1s 22ms/step - loss: 1.7468 - accuracy: 0.1989 - val\_loss: 1.8834 - val\_accuracy: 0.1916  
Epoch 24/50  
39/39 [=====] - 1s 26ms/step - loss: 1.7362 - accuracy: 0.2035 - val\_loss: 1.9160 - val\_accuracy: 0.1724  
Epoch 25/50  
39/39 [=====] - 1s 26ms/step - loss: 1.7442 - accuracy: 0.1991 - val\_loss: 1.9079 - val\_accuracy: 0.1762

```
Epoch 26/50
39/39 [=====] - 1s 22ms/step - loss: 1.7374 - accuracy: 0.2078 - val_loss: 1.9109 - val_accuracy: 0.1801
Epoch 27/50
39/39 [=====] - 1s 21ms/step - loss: 1.7371 - accuracy: 0.2049 - val_loss: 1.8901 - val_accuracy: 0.1762
Epoch 28/50
39/39 [=====] - 1s 22ms/step - loss: 1.7401 - accuracy: 0.2049 - val_loss: 1.9109 - val_accuracy: 0.1916
Epoch 29/50
```

## Evaluating Model

```
y_pred = model.predict(test_mfcc)
print("\n")
print("Accuracy score =",accuracy_score(y_test.argmax(axis=1),y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =",f1_score(y_test.argmax(axis=1),y_pred.argmax(axis=1),average = 'macro'))
print("\n")
print("F1 score (weighted average) =",f1_score(y_test.argmax(axis=1),y_pred.argmax(axis=1),average = 'weighted'))
```


```
70/70 [=====] - 0s 4ms/step
```

```
Accuracy score = 0.17107030900134348
```

```
F1 score (macro average) = 0.04869343530911408
```

```
F1 score (weighted average) = 0.04998000614800245
```

```
con_mat = confusion_matrix(y_test.argmax(axis=1),y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                 xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                 yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```

class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)
#print(precisions)
#print(recalls)

```

```

class_idx = list(range(len(class_labels)))
sorted_classes = sorted(class_idx, key=lambda x: f1_scores[x])

k=1
print("Most Confusing Classes:\n")
for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

```

1 - ANG:
    F1-score = nan    Precision = nan    Recall = 0.0000
2 - DIS:
    F1-score = nan    Precision = nan    Recall = 0.0000
3 - FEA:
    F1-score = nan    Precision = nan    Recall = 0.0000
4 - HAP:
    F1-score = nan    Precision = nan    Recall = 0.0000
5 - NEU:
    F1-score = nan    Precision = nan    Recall = 0.0000
6 - SAD:
    F1-score = 0.2922    Precision = 0.1711    Recall = 1.0000
/tmp/ipykernel_34/2911065269.py:13: RuntimeWarning: invalid value encountered in long_scalars
precision = TP / (TP + FP)

```

## Model 1D - Energy, Zero Crossing Rate and MFCCs

### Building Model

```

model = Sequential()
model.add(Conv1D(16, 12, activation='relu', input_shape=((np.concatenate((train_mfcc,train_energy,train_zcr), axis=1)).shape[1], 1)))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 8, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Conv1D(24, 5, activation='relu'))
model.add(MaxPooling1D(pool_size=2, strides=2))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
#model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))

```

```

model.add(Dense(6, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.001),
              metrics=['accuracy'])

model.summary()

```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
=====		
conv1d_27 (Conv1D)	(None, 460, 16)	208
max_pooling1d_27 (MaxPooling1D)	(None, 230, 16)	0
conv1d_28 (Conv1D)	(None, 223, 24)	3096
max_pooling1d_28 (MaxPooling1D)	(None, 111, 24)	0
conv1d_29 (Conv1D)	(None, 107, 24)	2904
max_pooling1d_29 (MaxPooling1D)	(None, 53, 24)	0
flatten_9 (Flatten)	(None, 1272)	0
dense_18 (Dense)	(None, 64)	81472
dropout_9 (Dropout)	(None, 64)	0
dense_19 (Dense)	(None, 6)	390
=====		
Total params: 88,070		
Trainable params: 88,070		
Non-trainable params: 0		
=====		

```

history = model.fit((np.concatenate((train_mfcc, train_energy, train_zcr), axis=1)), y_train,
                    epochs=50,
                    batch_size=128,
                    validation_data=((np.concatenate((val_mfcc, val_energy, val_zcr), axis=1)), y_val))

```



Epoch 1/50  
39/39 [=====] - 3s 53ms/step - loss: 2.0548 - accuracy: 0.1975 - val\_loss: 1.7181 - val\_accuracy: 0.2261  
Epoch 2/50  
39/39 [=====] - 2s 44ms/step - loss: 1.6412 - accuracy: 0.3145 - val\_loss: 1.5258 - val\_accuracy: 0.3372  
Epoch 3/50  
39/39 [=====] - 2s 45ms/step - loss: 1.5426 - accuracy: 0.3575 - val\_loss: 1.4569 - val\_accuracy: 0.3908  
Epoch 4/50  
39/39 [=====] - 2s 45ms/step - loss: 1.5170 - accuracy: 0.3650 - val\_loss: 1.3878 - val\_accuracy: 0.4483  
Epoch 5/50  
39/39 [=====] - 2s 44ms/step - loss: 1.4805 - accuracy: 0.3765 - val\_loss: 1.3771 - val\_accuracy: 0.4444  
Epoch 6/50  
39/39 [=====] - 2s 44ms/step - loss: 1.4726 - accuracy: 0.3943 - val\_loss: 1.3645 - val\_accuracy: 0.4444  
Epoch 7/50  
39/39 [=====] - 2s 54ms/step - loss: 1.4569 - accuracy: 0.3983 - val\_loss: 1.3597 - val\_accuracy: 0.4636  
Epoch 8/50  
39/39 [=====] - 2s 50ms/step - loss: 1.4441 - accuracy: 0.4072 - val\_loss: 1.3612 - val\_accuracy: 0.4444  
Epoch 9/50  
39/39 [=====] - 2s 45ms/step - loss: 1.4336 - accuracy: 0.4167 - val\_loss: 1.3412 - val\_accuracy: 0.4483  
Epoch 10/50  
39/39 [=====] - 2s 47ms/step - loss: 1.4296 - accuracy: 0.4091 - val\_loss: 1.3414 - val\_accuracy: 0.4636  
Epoch 11/50  
39/39 [=====] - 2s 43ms/step - loss: 1.4200 - accuracy: 0.4254 - val\_loss: 1.3283 - val\_accuracy: 0.4483  
Epoch 12/50  
39/39 [=====] - 2s 43ms/step - loss: 1.4096 - accuracy: 0.4258 - val\_loss: 1.3198 - val\_accuracy: 0.4674  
Epoch 13/50  
39/39 [=====] - 2s 42ms/step - loss: 1.4046 - accuracy: 0.4337 - val\_loss: 1.3268 - val\_accuracy: 0.4674  
Epoch 14/50  
39/39 [=====] - 2s 42ms/step - loss: 1.3919 - accuracy: 0.4337 - val\_loss: 1.2956 - val\_accuracy: 0.4828  
Epoch 15/50  
39/39 [=====] - 2s 43ms/step - loss: 1.3860 - accuracy: 0.4406 - val\_loss: 1.3141 - val\_accuracy: 0.4789  
Epoch 16/50  
39/39 [=====] - 2s 44ms/step - loss: 1.3970 - accuracy: 0.4305 - val\_loss: 1.3176 - val\_accuracy: 0.4636  
Epoch 17/50  
39/39 [=====] - 2s 42ms/step - loss: 1.3800 - accuracy: 0.4450 - val\_loss: 1.2926 - val\_accuracy: 0.4713  
Epoch 18/50  
39/39 [=====] - 2s 42ms/step - loss: 1.3821 - accuracy: 0.4414 - val\_loss: 1.2904 - val\_accuracy: 0.5019  
Epoch 19/50  
39/39 [=====] - 2s 43ms/step - loss: 1.3760 - accuracy: 0.4416 - val\_loss: 1.2972 - val\_accuracy: 0.4751  
Epoch 20/50  
39/39 [=====] - 2s 44ms/step - loss: 1.3670 - accuracy: 0.4434 - val\_loss: 1.3036 - val\_accuracy: 0.4828  
Epoch 21/50  
39/39 [=====] - 2s 44ms/step - loss: 1.3636 - accuracy: 0.4503 - val\_loss: 1.2929 - val\_accuracy: 0.4904  
Epoch 22/50  
39/39 [=====] - 2s 46ms/step - loss: 1.3523 - accuracy: 0.4630 - val\_loss: 1.2739 - val\_accuracy: 0.5019  
Epoch 23/50  
39/39 [=====] - 2s 47ms/step - loss: 1.3545 - accuracy: 0.4580 - val\_loss: 1.2820 - val\_accuracy: 0.4904  
Epoch 24/50  
39/39 [=====] - 2s 44ms/step - loss: 1.3710 - accuracy: 0.4531 - val\_loss: 1.2895 - val\_accuracy: 0.4828

```
Epoch 25/50
39/39 [=====] - 2s 58ms/step - loss: 1.3356 - accuracy: 0.4592 - val_loss: 1.2969 - val_accuracy: 0.4751
Epoch 26/50
39/39 [=====] - 2s 47ms/step - loss: 1.3388 - accuracy: 0.4578 - val_loss: 1.2878 - val_accuracy: 0.4904
Epoch 27/50
39/39 [=====] - 2s 48ms/step - loss: 1.3292 - accuracy: 0.4685 - val_loss: 1.2722 - val_accuracy: 0.4789
Epoch 28/50
39/39 [=====] - 2s 50ms/step - loss: 1.3350 - accuracy: 0.4648 - val_loss: 1.2827 - val_accuracy: 0.4789
Epoch 29/50
39/39 [=====] - 2s 48ms/step - loss: 1.3268 - accuracy: 0.4671 - val_loss: 1.2684 - val_accuracy: 0.4943
```

## Evaluating Model

```
y_pred = model.predict((np.concatenate((test_mfcc, test_energy, test_zcr), axis=1)))
print("\n")
print("Accuracy score =", accuracy_score(y_test.argmax(axis=1), y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'macro'))
print("\n")
print("F1 score (weighted average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'weighted'))
```

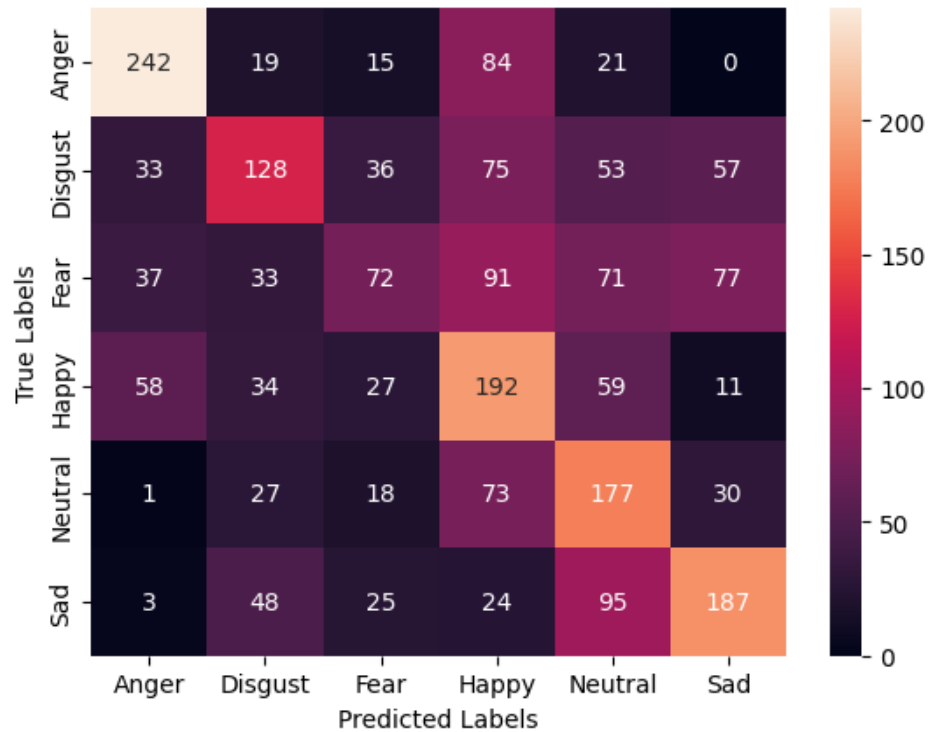
```
70/70 [=====] - 1s 6ms/step
```

```
Accuracy score = 0.44693237796686075
```

```
F1 score (macro average) = 0.4391544645220722
```

```
F1 score (weighted average) = 0.4391018836581617
```

```
con_mat = confusion_matrix(y_test.argmax(axis=1), y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                  xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                  yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```

class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)

```

```

#print(precisions)
#print(recalls)

class_idx = list(range(len(class_labels)))
sorted_classes = sorted(class_idx, key=lambda x: f1_scores[x])

k=1
print("Most Confusing Classes:\n")
for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

1 - FEA:	F1-score = 0.2509	Precision = 0.3731	Recall = 0.1890
2 - DIS:	F1-score = 0.3815	Precision = 0.4429	Recall = 0.3351
3 - HAP:	F1-score = 0.4174	Precision = 0.3562	Recall = 0.5039
4 - NEU:	F1-score = 0.4414	Precision = 0.3718	Recall = 0.5429
5 - SAD:	F1-score = 0.5027	Precision = 0.5166	Recall = 0.4895
6 - ANG:	F1-score = 0.6411	Precision = 0.6471	Recall = 0.6352

## Model 2D - Mel Spectrogram

### Building Model

We build our own model but we have noticed that the accuracy not so good as shown.

```

model = Sequential()
model.add(Conv2D(16, (7, 7), activation='relu', input_shape=(train_mel.shape[1], train_mel.shape[2], 1)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Conv2D(24, (7, 7), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

```

```

model.add(BatchNormalization())
model.add(Conv2D(24, (7, 7), activation='relu'))
model.add(Flatten())
model.add(Dense(64, activation='relu', kernel_regularizer=regularizers.l2(0.00001)))
model.add(BatchNormalization())
model.add(Dropout(0.2))
model.add(Dense(6, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.00001),
              metrics=['accuracy'])

model.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 122, 151, 16)	800
max_pooling2d_2 (MaxPooling 2D)	(None, 61, 75, 16)	0
batch_normalization_3 (Batch Normalization)	(None, 61, 75, 16)	64
conv2d_4 (Conv2D)	(None, 55, 69, 24)	18840
max_pooling2d_3 (MaxPooling 2D)	(None, 27, 34, 24)	0
batch_normalization_4 (Batch Normalization)	(None, 27, 34, 24)	96
conv2d_5 (Conv2D)	(None, 21, 28, 24)	28248
flatten_1 (Flatten)	(None, 14112)	0
dense_2 (Dense)	(None, 64)	903232
batch_normalization_5 (Batch Normalization)	(None, 64)	256
dropout_1 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 6)	390

```
=====
Total params: 951,926
Trainable params: 951,718
Non-trainable params: 208

```

---

```
history = model.fit(train_mel, y_train,
                    epochs=10,
                    batch_size=128,
                    validation_data=(val_mel, y_val))
```

```
Epoch 1/10
39/39 [=====] - 229s 6s/step - loss: 2.3071 - accuracy: 0.1964 - val_loss: 2.3234 - val_accuracy: 0.2031
Epoch 2/10
39/39 [=====] - 222s 6s/step - loss: 1.9433 - accuracy: 0.2920 - val_loss: 2.2687 - val_accuracy: 0.2069
Epoch 3/10
39/39 [=====] - 223s 6s/step - loss: 1.7613 - accuracy: 0.3472 - val_loss: 2.1264 - val_accuracy: 0.2222
Epoch 4/10
39/39 [=====] - 223s 6s/step - loss: 1.6761 - accuracy: 0.3707 - val_loss: 1.9395 - val_accuracy: 0.2529
Epoch 5/10
39/39 [=====] - 223s 6s/step - loss: 1.6030 - accuracy: 0.3941 - val_loss: 1.7752 - val_accuracy: 0.3103
Epoch 6/10
39/39 [=====] - 221s 6s/step - loss: 1.5754 - accuracy: 0.3957 - val_loss: 1.6510 - val_accuracy: 0.3602
Epoch 7/10
39/39 [=====] - 226s 6s/step - loss: 1.5214 - accuracy: 0.4194 - val_loss: 1.6229 - val_accuracy: 0.3870
Epoch 8/10
39/39 [=====] - 222s 6s/step - loss: 1.4913 - accuracy: 0.4321 - val_loss: 1.5827 - val_accuracy: 0.4100
Epoch 9/10
39/39 [=====] - 224s 6s/step - loss: 1.4522 - accuracy: 0.4440 - val_loss: 1.5170 - val_accuracy: 0.3946
Epoch 10/10
39/39 [=====] - 222s 6s/step - loss: 1.4486 - accuracy: 0.4444 - val_loss: 1.5055 - val_accuracy: 0.4100
```

## Evaluating Model

```
y_pred = model.predict(test_mel)
print("\n")
print("Accuracy score =", accuracy_score(y_test.argmax(axis=1), y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'macro'))
print("\n")
print("F1 score (weighted average) =", f1_score(y_test.argmax(axis=1), y_pred.argmax(axis=1), average = 'weighted'))
```

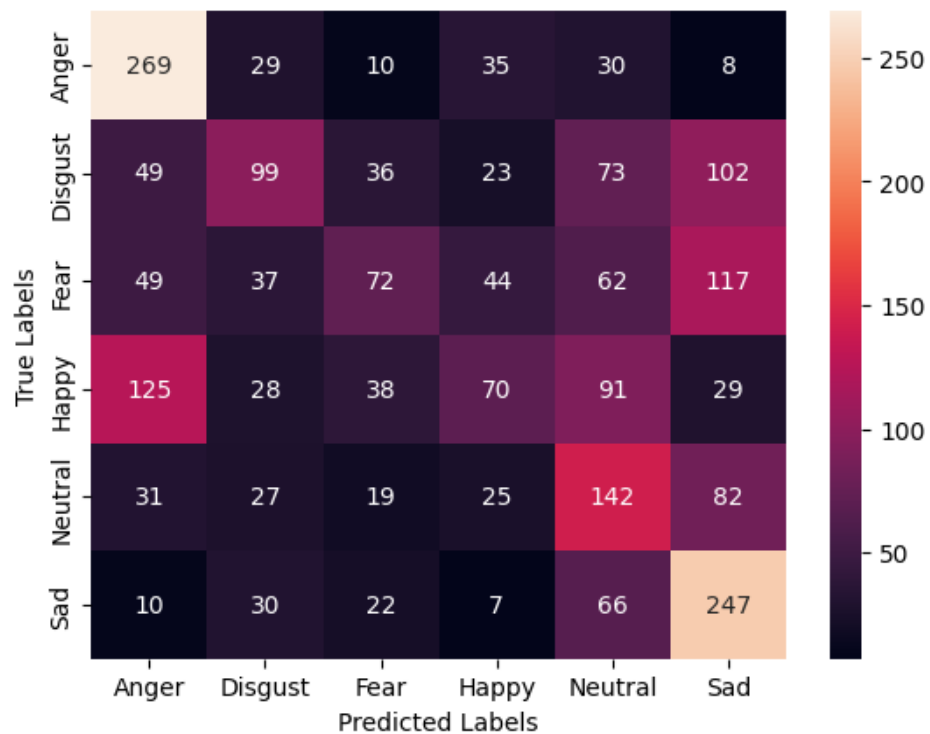
```
70/70 [=====] - 26s 365ms/step
```

Accuracy score = 0.4025974025974026

F1 score (macro average) = 0.37678596107610596

F1 score (weighted average) = 0.3772434861554859

```
con_mat = confusion_matrix(y_test.argmax(axis=1),y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                 xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                 yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```

class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)
#print(precisions)
#print(recalls)

class_idxes = list(range(len(class_labels)))
sorted_classes = sorted(class_idxes, key=lambda x: f1_scores[x])

k=1
print("Most Confusing Classes:\n")
for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

1 - FEA:	F1-score = 0.2665	Precision = 0.2934	Recall = 0.2441
2 - DIS:	F1-score = 0.2778	Precision = 0.3696	Recall = 0.2225
3 - NEU:	F1-score = 0.3156	Precision = 0.4150	Recall = 0.2546
4 - HAP:	F1-score = 0.3703	Precision = 0.3679	Recall = 0.3727
5 - SAD:			



	F1-score = 0.4513	Precision = 0.4422	Recall = 0.4607
6 - ANG:	F1-score = 0.5485	Precision = 0.4231	Recall = 0.7795

## Model 2D - Research Paper

This model is replicated from a model in this research paper with some modifications and it gives higher accuracy and less confusing classes.

### **Emotion Recognition in Audio and Video Using Deep Neural Networks**

Mandeep Singh  
SCPD, Stanford University  
Stanford, CA  
[msingh13@stanford.edu](mailto:msingh13@stanford.edu)

Yuan Fang  
ICME, Stanford University  
Stanford, CA  
[yuanfy@stanford.edu](mailto:yuanfy@stanford.edu)

## Building Model

This was the CNN model architecture used in the original research paper, we have replicated it here. However, since the research paper was used on a different dataset, we have modified the input and output layers accordingly.

# CNN

INPUT	200x300
CONV 1	16 filters of 12x16
ReLU	
MaxPool2D	Size 2 with Stride 2
CONV 2	24 filters of 8x12
ReLU	
MaxPool2D	Size 2 with Stride 2
CONV 3	24 filters of 5x7
ReLU	
MaxPool2D	Size 2 with Stride 2
Flatten	
Linear	64
ReLU	
Dropout	0.2
Linear	4

```
model = Sequential()
model.add(Conv2D(16, (12, 6), activation='relu', input_shape=(train_mel.shape[1], train_mel.shape[2], 1)))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Conv2D(24, (8, 12), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Conv2D(24, (5, 7), activation='relu'))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(6, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(learning_rate=0.001),
              metrics=['accuracy'])

model.summary()
```

Model: "sequential\_18"

---

Layer (type)	Output Shape	Param #
conv2d_21 (Conv2D)	(None, 117, 152, 16)	1168
max_pooling2d_16 (MaxPooling2D)	(None, 58, 76, 16)	0
conv2d_22 (Conv2D)	(None, 51, 65, 24)	36888
max_pooling2d_17 (MaxPooling2D)	(None, 25, 32, 24)	0
conv2d_23 (Conv2D)	(None, 21, 26, 24)	20184
flatten_18 (Flatten)	(None, 13104)	0
dense_36 (Dense)	(None, 64)	838720
dropout_18 (Dropout)	(None, 64)	0
dense_37 (Dense)	(None, 6)	390
Total params: 897,350		
Trainable params: 897,350		
Non-trainable params: 0		

```
history = model.fit(train_mel, y_train,
                    epochs=10,
                    batch_size=128,
                    validation_data=(val_mel, y_val))
```

```
Epoch 1/10
39/39 [=====] - 155s 4s/step - loss: 3.2869 - accuracy: 0.1776 - val_loss: 1.7848 - val_accuracy: 0.2184
Epoch 2/10
39/39 [=====] - 154s 4s/step - loss: 1.7462 - accuracy: 0.2478 - val_loss: 1.5522 - val_accuracy: 0.3333
Epoch 3/10
39/39 [=====] - 154s 4s/step - loss: 1.5277 - accuracy: 0.3832 - val_loss: 1.3663 - val_accuracy: 0.4483
Epoch 4/10
39/39 [=====] - 154s 4s/step - loss: 1.3874 - accuracy: 0.4464 - val_loss: 1.3026 - val_accuracy: 0.4598
Epoch 5/10
39/39 [=====] - 153s 4s/step - loss: 1.3019 - accuracy: 0.4893 - val_loss: 1.3870 - val_accuracy: 0.4521
Epoch 6/10
39/39 [=====] - 154s 4s/step - loss: 1.2287 - accuracy: 0.5253 - val_loss: 1.2957 - val_accuracy: 0.4444
Epoch 7/10
39/39 [=====] - 156s 4s/step - loss: 1.1668 - accuracy: 0.5451 - val_loss: 1.2497 - val_accuracy: 0.4713
```

```
Epoch 8/10
39/39 [=====] - 154s 4s/step - loss: 1.0673 - accuracy: 0.5835 - val_loss: 1.2772 - val_accuracy: 0.4789
Epoch 9/10
39/39 [=====] - 153s 4s/step - loss: 0.9423 - accuracy: 0.6372 - val_loss: 1.3159 - val_accuracy: 0.4751
Epoch 10/10
39/39 [=====] - 155s 4s/step - loss: 0.8536 - accuracy: 0.6688 - val_loss: 1.3401 - val_accuracy: 0.5096
```

## Evaluating Model

```
y_pred = model.predict(test_mel)
print("\n")
print("Accuracy score =",accuracy_score(y_test.argmax(axis=1),y_pred.argmax(axis=1)))
print("\n")
print("F1 score (macro average) =",f1_score(y_test.argmax(axis=1),y_pred.argmax(axis=1),average = 'macro'))
print("\n")
print("F1 score (weighted average) =",f1_score(y_test.argmax(axis=1),y_pred.argmax(axis=1),average = 'weighted'))
```


```
70/70 [=====] - 15s 208ms/step
```

```
Accuracy score = 0.46618898343036275
```

```
F1 score (macro average) = 0.4614819803410712
```

```
F1 score (weighted average) = 0.46092974621199106
```

```
con_mat = confusion_matrix(y_test.argmax(axis=1),y_pred.argmax(axis=1))
#unique_labels = np.unique(labels)
#print(unique_labels)
ax = sns.heatmap(con_mat, annot=True, fmt="d",
                 xticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'],
                 yticklabels = ['Anger', 'Disgust', 'Fear', 'Happy', 'Neutral', 'Sad'])
#ax = sns.heatmap(con_mat, annot=True, fmt="d", xticklabels=unique_labels, yticklabels=unique_labels)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```

class_labels = list(np.unique(labels))
#print(class_labels)

precisions = []
recalls = []
f1_scores = []

for i in range(len(class_labels)):
    TP = con_mat[i, i]
    FP = np.sum(con_mat[:, i]) - TP
    FN = np.sum(con_mat[i, :]) - TP

    precision = TP / (TP + FP)
    recall = TP / (TP + FN)
    f1 = 2 * (precision * recall) / (precision + recall)

    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#print(f1_scores)
#print(precisions)
#print(recalls)

```

```

class_idx = list(range(len(class_labels)))
sorted_classes = sorted(class_idx, key=lambda x: f1_scores[x])

k=1
print("Most Confusing Classes:\n")
for j in sorted_classes:
    label = class_labels[j]
    print(f"{k} - {label}:\n\tF1-score = {f1_scores[j]:.4f}\tPrecision = {precisions[j]:.4f}\tRecall = {recalls[j]:.4f} ")
    k+=1

```

Most Confusing Classes:

```

1 - DIS:
    F1-score = 0.3381      Precision = 0.4327      Recall = 0.2775
2 - FEA:
    F1-score = 0.3793      Precision = 0.3834      Recall = 0.3753
3 - HAP:
    F1-score = 0.4470      Precision = 0.4307      Recall = 0.4646
4 - NEU:
    F1-score = 0.4826      Precision = 0.4586      Recall = 0.5092
5 - SAD:
    F1-score = 0.5109      Precision = 0.4752      Recall = 0.5524
6 - ANG:
    F1-score = 0.6110      Precision = 0.5980      Recall = 0.6247

```

These were the results obtained in the original paper, note that the datasets used were different, so labels and accuracies will also be slightly different

Architecture	Accuracy(%)	Data Aug.	Emotion
CNN	52.23	No	H,S,A,N
CNN	51.90	Yes	H,S,A,N

